





# Regular expressions

- A formal language for specifying text strings
- How can we search for any of these?
  - woodchuck
  - woodchucks
  - Woodchuck
  - Woodchucks





# Regular Expressions: Disjunctions

- Letters inside square brackets []

| Pattern      | Matches              |
|--------------|----------------------|
| [wW]oodchuck | Woodchuck, woodchuck |
| [1234567890] | Any digit            |

- Ranges [A-Z]

| Pattern | Matches              |   |
|---------|----------------------|---|
| [A-Z]   | An upper case letter | <u>D</u> renched Blossoms               |
| [a-z]   | A lower case letter  | <u>m</u> y beans were impatient         |
| [0-9]   | A single digit       | Chapter <u>1</u> : Down the Rabbit Hole |



# Regular Expressions: Negation in Disjunction

- Negations [ ^Ss ]
  - Carat means negation only when first in []

| Pattern  | Matches                  |                                    |
|----------|--------------------------|------------------------------------|
| [ ^A-Z ] | Not an upper case letter | Oyfn pripetchik                    |
| [ ^Ss ]  | Neither 'S' nor 's'      | <u>I</u> have no exquisite reason" |
| [ ^e^ ]  | Neither e nor ^          | Look h <u>e</u> re                 |
| a^b      | The pattern a carat b    | Look up <u>a^b</u> now             |



# Regular Expressions: More Disjunction

- Woodchucks is another name for groundhog!
- The pipe | for disjunction


| Pattern                                | Matches              |
|--|----------------------|
| <code>groundhog woodchuck</code>       |                      |
| <code>yours mine</code>                | yours<br>mine        |
| <code>a b c</code>                     | = <code>[abc]</code> |
| <code>[gG]roundhog [Ww]oodchuck</code> |                      |





# Regular Expressions: ? \* + .

| Pattern              | Matches                    |   |
|----------------------|----------------------------|---|
| <code>colou?r</code> | Optional previous char     | <u>color</u> <u>colour</u>                          |
| <code>oo*h!</code>   | 0 or more of previous char | <u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>   |
| <code>o+h!</code>    | 1 or more of previous char | <u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>   |
| <code>baa+</code>    |                            | <u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>   |
| <code>beg.n</code>   |                            | <u>begin</u> <u>begun</u> <u>begun</u> <u>beg3n</u> |



Stephen C Kleene

Kleene \*, Kleene +



# Regular Expressions: Anchors <sup>^</sup> \$

| Pattern                             | Matches                           |
|-------------------------------------|-----------------------------------|
| <sup>^</sup> [A-Z]                  | <u>P</u> alo Alto                 |
| <sup>^</sup> [ <sup>^</sup> A-Za-z] | <u>1</u> <u>"Hello"</u>           |
| \. \$                               | The end <u>.</u>                  |
| . \$                                | The end <u>?</u> The end <u>!</u> |



## Example

- Find me all instances of the word “the” in a text.

the

Misses capitalized examples

[tT]he

Incorrectly returns other or theology

[^a-zA-Z][tT]he[^a-zA-Z]





## Errors

- The process we just went through was based on **fixing two kinds of errors**
  - Matching strings that we should not have matched (**there**, **then**, **other**)
    - **False positives (Type I)**
  - Not matching things that we should have matched (**The**)
    - **False negatives (Type II)**



## Errors cont.

- In NLP we are always dealing with these kinds of errors.
- Reducing the error rate for an application often involves two antagonistic efforts:
  - Increasing accuracy or precision (minimizing false positives)
  - Increasing coverage or recall (minimizing false negatives).



## Summary

- Regular expressions play a surprisingly large role
  - Sophisticated sequences of regular expressions are often the first model for any text processing text
- For many hard tasks, we use machine learning classifiers
  - But regular expressions are used as features in the classifiers
  - Can be very useful in capturing generalizations



# Word tokenization





## Text Normalization

- Every NLP task needs to do text normalization:
  1. Segmenting/tokenizing words in running text
  2. Normalizing word formats
  3. Segmenting sentences in running text



## How many words?

- I do uh main- mainly business data processing
  - Fragments, filled pauses
- Seuss's **cat** in the hat is different from other **cats**!
  - **Lemma**: same stem, part of speech, rough word sense
    - **cat** and **cats** = same lemma    Lemma : 代表相同的字，只是變化型不一樣
  - **Wordform**: the full inflected surface form
    - **cat** and **cats** = different wordforms



## How many words?

they lay back on the San Francisco grass and looked at the stars and their

- **Type:** an element of the vocabulary. 用Lemma角度去計算
- **Token:** an instance of that type in running text. 用wordform角度去計算
- How many?
  - 15 tokens (or 14)
  - 13 types (or 12) (or 11?)  
and, and  
they, their  
the, the





# How many words?

**$N$**  = number of tokens

**$V$**  = vocabulary = set of types

$|V|$  is the size of the vocabulary

Church and Gale (1990):  $|V| > O(N^{1/2})$

Wordform

Lemma

|                                 | Tokens = $N$ | Types = $ V $ |
|---------------------------------|--------------|---------------|
| Switchboard phone conversations | 2.4 million  | 20 thousand   |
| Shakespeare                     | 884,000      | 31 thousand   |
| Google N-grams                  | 1 trillion   | 13 million    |



# Simple Tokenization in UNIX

- (Inspired by Ken Church's UNIX for Poets.)
- Given a text file, output the word tokens and their frequencies

```
tr -sc 'A-Za-z' '\n' < shakes.txt
```

Change all non-alpha to newlines

```
| sort
```

Sort in alphabetical order

```
| uniq -c
```

Merge and count each type

|           |          |
|-----------|----------|
| 1945 A    | 25 Aaron |
| 72 AARON  | 6 Abate  |
| 19 ABBESS | 1 Abates |
| 5 ABBOT   | 5 Abbess |
| ...       | 6 Abbey  |
| ...       | 3 Abbot  |
| ...       | ....     |
| ...       | ...      |

Dan Jurafsky



## The first step: tokenizing

```
tr -sc 'A-Za-z' '\n' < shakes.txt | head
```

```
THE  
SONNETS  
by  
William  
Shakespeare  
From  
fairest  
creatures  
We  
...
```

Dan Jurafsky



## The second step: sorting

```
tr -sc 'A-Za-z' '\n' < shakes.txt | sort | head
```

A

A

A

A

A

A

A

A

A

...



## More counting

- Merging upper and lower case

```
tr 'A-Z' 'a-z' < shakes.txt | tr -sc 'A-Za-z' '\n' | sort | uniq -c
```

- Sorting the counts

```
tr 'A-Z' 'a-z' < shakes.txt | tr -sc 'A-Za-z' '\n' | sort | uniq -c | sort -n -r
```

```
23243 the
22225 i
18618 and
16339 to
15687 of
12780 a
12163 you
10839 my
10005 in
8954 d
```

What happened here?



## Issues in Tokenization

- Finland's capital → Finland Finlands Finland's ?
- what're, I'm, isn't → What are, I am, is not
- Hewlett-Packard → Hewlett Packard ?
- state-of-the-art → state of the art ?
- Lowercase → lower-case lowercase lower case ?
- San Francisco → one token or two?
- m.p.h., PhD. → ??



# Tokenization: language issues

- French
  - *L'ensemble* → one token or two?
    - *L ? L' ? Le ?*
    - Want *l'ensemble* to match with *un ensemble*
- German noun compounds are not segmented
  - *Lebensversicherungsgesellschaftsangestellter*
  - 'life insurance company employee'
  - German information retrieval needs **compound splitter**



## Tokenization: language issues

- Chinese and Japanese no spaces between words:
  - 莎拉波娃现在居住在美国东南部的佛罗里达。
  - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
  - Sharapova now lives in US southeastern Florida
- Further complicated in Japanese, with multiple alphabets intermingled
  - Dates/amounts in multiple formats

フォーチュン500社は情報不足のため時間あた\$500K(約6,000万円)

Katakana      Hiragana      Kanji      Romaji

End-user can express query entirely in hiragana!





# Word Tokenization in Chinese

- Also called **Word Segmentation**
- Chinese words are composed of characters
  - Characters are generally 1 syllable and 1 morpheme.
  - Average word is 2.4 characters long.
- Standard baseline segmentation algorithm:
  - **Maximum Matching** (also called Greedy)



# Maximum Matching Word Segmentation Algorithm

- Given a wordlist of Chinese, and a string.
  - 1) Start a pointer at the beginning of the string
  - 2) Find the longest word in dictionary that matches the string starting at pointer
  - 3) Move the pointer over the word in string
  - 4) Go to 2

從輸入的字串當中

找出最長和字典中的字彙相同的字

即可做斷詞

針對句子本身沒有空格的語言比較有效果（中文、日文）

Dan Jurafsky



## Max-match segmentation illustration

- Thecatinthehat                      the cat in the hat
- Thetabledownthere                the table down there  
    theta bled own there
- Doesn't generally work in English!
- But works astonishingly well in Chinese
  - 莎拉波娃现在居住在美国东南部的佛罗里达。
  - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
- Modern probabilistic segmentation algorithms even better

# Word tokenization



[illegible][illegible]



# Normalization

要把所有字的變化型轉換為一致的型態  
這樣在後續的處理上才能比較符合原意

- Need to “normalize” terms
  - Information Retrieval: indexed text & query terms must have same form.
    - We want to match ***U.S.A.*** and ***USA***
- We implicitly define equivalence classes of terms
  - e.g., deleting periods in a term
- Alternative: asymmetric expansion:
  - Enter: ***window***                      Search: ***window, windows***
  - Enter: ***windows***                      Search: ***Windows, windows, window***
  - Enter: ***Windows***                      Search: ***Windows***
- Potentially more powerful, but less efficient



## Case folding

- Applications like IR: reduce all letters to lower case
  - Since users tend to use lower case
  - Possible exception: upper case in mid-sentence?
    - e.g., *General Motors*
    - *Fed* vs. *fed*
    - *SAIL* vs. *sail*
- For sentiment analysis, MT, Information extraction
  - Case is helpful (*US* versus *us* is important)



# Lemmatization

把所有變化型轉為原型

- Reduce inflections or variant forms to base form
  - *am, are, is* → *be*
  - *car, cars, car's, cars'* → *car*
- *the boy's cars are different colors* → *the boy car be different color*
- Lemmatization: have to find correct dictionary headword form
- Machine translation
  - Spanish **quiero** ('I want'), **quieres** ('you want') same lemma as **querer** 'want'





# Morphology

- **Morphemes:**
  - The small meaningful units that make up words
  - **Stems:** The core meaning-bearing units
  - **Affixes:** Bits and pieces that adhere to stems
    - Often with grammatical functions

Stem: 一個字的核心部分

Affixes: 字的變化型部分

e.g. affordable

Stem: afford

Affixes: able



# Stemming

- Reduce terms to their stems in information retrieval
- *Stemming* is crude chopping of affixes
  - language dependent
  - e.g., ***automate(s), automatic, automation*** all reduced to ***automat***.

*for example compressed  
and compression are both  
accepted as equivalent to  
compress.*



for exampl compress and  
compress ar both accept  
as equal to compress



# Porter's algorithm

## The most common English stemmer

### Step 1a

|           |                   |
|-----------|-------------------|
| sses → ss | caresses → caress |
| ies → i   | ponies → poni     |
| ss → ss   | caress → caress   |
| s → ∅     | cats → cat        |

### Step 1b

|              |                     |
|--------------|---------------------|
| (*v*)ing → ∅ | walking → walk      |
|              | sing → sing         |
| (*v*)ed → ∅  | plastered → plaster |
| 母音+ing/ed    |                     |
| ...          |                     |

### Step 2 (for long stems)

|               |                      |
|---------------|----------------------|
| ational → ate | relational → relate  |
| izer → ize    | digitizer → digitize |
| ator → ate    | operator → operate   |
| ...           |                      |

### Step 3 (for longer stems)

|          |                     |
|----------|---------------------|
| al → ∅   | revival → reviv     |
| able → ∅ | adjustable → adjust |
| ate → ∅  | activate → activ    |
| ...      |                     |

# Viewing morphology in a corpus

( \*v\* )ing  $\rightarrow$   $\emptyset$  walking  $\rightarrow$  walk

sing → sing

(*\*v\**)ing → ∅    walking → walk  
                     sing → sing

|      |               |     |               |
|------|---------------|-----|---------------|
| 1312 | King          | 548 | being         |
| 548  | <b>being</b>  | 541 | nothing       |
| 541  | nothing       | 152 | something     |
| 388  | king          | 145 | <b>coming</b> |
| 375  | bring         | 130 | morning       |
| 358  | thing         | 122 | having        |
| 307  | ring          | 120 | living        |
| 152  | something     | 117 | loving        |
| 145  | <b>coming</b> | 116 | Being         |
| 130  | morning       | 102 | going         |

37



# Dealing with complex morphology is sometimes necessary

- Some languages requires complex morpheme segmentation
  - Turkish
  - **Uygarlastiramadiklarimizdanmissinizcasina**
  - `(behaving) as if you are among those whom we could not civilize’
  - **Uygar** `civilized’ + **las** `become’
    - + **tir** `cause’ + **ama** `not able’
    - + **dik** `past’ + **lar** `plural’
    - + **imiz** `p1pl’ + **dan** `abl’
    - + **mis** `past’ + **siniz** `2pl’ + **casina** `as if’

[illegible]



# Basic Text Processing

# Sentence Segmentation and Decision Trees





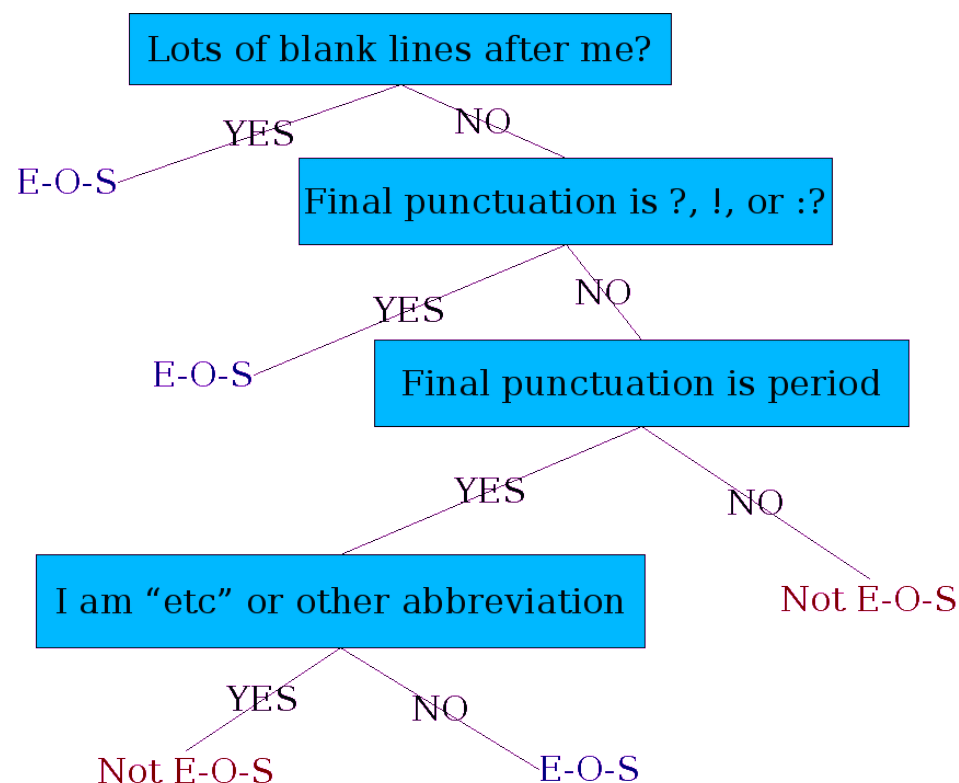
# Sentence Segmentation

- !, ? are relatively unambiguous
- Period “.” is quite ambiguous
  - Sentence boundary
  - Abbreviations like Inc. or Dr.
  - Numbers like .02% or 4.3
- Build a binary classifier
  - Looks at a “.”
  - Decides EndOfSentence/NotEndOfSentence
  - Classifiers: hand-written rules, regular expressions, or machine-learning

做句子切割，要如何找到句子的中斷點？光用標點符號判斷不一定準確  
要用分類器來判斷這個符號到底是不是End-of-sentence



# Determining if a word is end-of-sentence: a Decision Tree





## More sophisticated decision tree features

- Case of word with “.”: Upper, Lower, Cap, Number
- Case of word after “.”: Upper, Lower, Cap, Number
- Numeric features
  - Length of word with “.”
  - Probability(word with “.” occurs at end-of-s)
  - Probability(word after “.” occurs at beginning-of-s)

e.g. 有一句話長這樣 ... .The ...

句點後面接了大寫The，

那這個句點是end-of-sentence的機率就很高



# Implementing Decision Trees

- A decision tree is just an if-then-else statement
- The interesting research is choosing the features
- Setting up the structure is often too hard to do by hand
  - Hand-building only possible for very simple features, domains
    - For numeric features, it's too hard to pick each threshold
  - Instead, structure usually learned by machine learning from a training corpus



## Decision Trees and other classifiers

- We can think of the questions in a decision tree
- As features that could be exploited by any kind of classifier
  - Logistic regression
  - SVM
  - Neural Nets
  - etc.



# Basic Text Processing

# Sentence Segmentation and Decision Trees